


The Delegation Blind Spot



Successful execution is not the same as evidence for the next product decision.

01 THE AUDIT

Does this log support this future decision?



Example: two optimal hotel bookings can look identical. They need not support the same investment in soundproofing or transport links.

02 THREE DISTINCT PIECES OF EVIDENCE

PRIMARY / PROSPECTIVELY FROZEN

0 / 36

comparisons resolved

4,800 API requests

GPT-5.4 Mini and Nano on shared synthetic tasks. They choose a utility maximizer in 73.19% and 55.69% of field tasks, respectively.

Conservative intervals remain unresolved. This does not prove structural information loss.

FOLLOW-UP / EXPLORATORY

3 / 9

resolved per model

2,400 additional calls

Report the largest supplied preference weight. Each model repeats only 12 distinct prompts. Every report is correct.

Matched observation counts, not matched costs. Known input extraction, not human preference discovery.

PARSER / EXPLORATORY REANALYSIS

7 / 9

comparisons resolved

720 shared field records

Read the supplied weights directly. A known identity channel removes report-calibration uncertainty. No new API calls.

Zero incorrect resolutions in this sample. The profile is synthetic, not validated against a customer.

Coverage is marginal, not across all decisions: primary and model-report intervals use 96%; the parser uses a known channel and 96% field intervals. Follow-ups are not corrected for selection after the primary results.

03 WHAT THE MATHEMATICS SAYS

A contrast, not a full profile.

p = proportions of declared intent classes
A = calibrated class-to-log channel
q = observed log distribution
d = independently specified product-value contrast

$q = Ap$ Target: $d^T p$

Global identification: for known A, the contrast is identified from every feasible q exactly when d lies in the row space of A.

With uncertain A and q, a joint-mass linear program returns attainable bounds and endpoint witness mixtures. A range crossing zero supports abstention.

04 STRUCTURAL LOSS ≠ FINITE PRECISION

Identified can still mean imprecise.

Controlled channel: $A(\eta) = (1 - \eta)11^T/4 + \eta I$
Exact width is zero whenever $\eta > 0$.

The graph plots 'Finite-sample median width' (teal line with circles) and 'Exact width' (orange line with squares) against 'Channel strength η (categorical spacing)'. The x-axis has values 0, 0.01, 0.03, 0.1, 0.3, and 1. The y-axis has values 0.00, 0.05, and 0.10. The 'Exact width' starts at 0.10 for η=0 and drops to 0.00 for η=0.01, remaining at 0.00 for all higher η. The 'Finite-sample median width' starts at 0.10 for η=0, remains at 0.10 for η=0.01 and 0.03, then decreases to approximately 0.075 at η=0.1, 0.04 at η=0.3, and 0.00 at η=1.

Channel strength η (categorical spacing)
Joint intervals; positive mixture; d = (.06, .02, -.04, -.02). 2,560 calibration samples per class; 5,120 field samples. 200 repetitions per point; widths condition on feasibility.

05 THE INTERFACE QUESTION

Keep the source of a claim visible.



Discuss: Which fields justify their privacy and elicitation cost? When should teams collect a direct outcome instead? Do witness mixtures help analysts interpret uncertainty?